

Report of Hand-Written Digits Classification Problem

Junzhou Fang
ECE, UIUC
Champaign, United States
junzhou5@illinois.edu

Abstract—This report is a summary of the hand-written digits classification problem. We included three methods: logistic regression with feature extraction, perceptron, CNN, and a simple algorithm, to solve the problem.

Index Terms—Binary classification, logistic regression, CNN, MNIST

I. INTRODUCTION

Professor Do's research team has collected a large amount of image hand-written data, but some of the image are blurred during the collection due to technological errors. The goal of this report is to classify the images into two categories: blurred (represent as 1) and non-blurred (represent as -1).

II. DATA PREPROCESSING

A. Data Description

There are three sets of images: train set, validation set and test set. The train set contains 40000 images, the validation set contains 10000 images and the test set contains 8000 images. Each image is in size of (28,28) and is in black and white. In addition to images, there are three text files containing labels of images in each set. We use 1 to represent unblurred images and -1 to represent blurred images.

B. Data Loading

We use skimage to load all the images and convert them into a numpy array. We use flatten() to construct image sets into a 2D array, with each image is a vector.

We also read labels into a numpy array. The index of the

```
import numpy as np
from skimage.io import imread
from skimage.transform import resize
import os

num_train = 39999
image_dir = 'cig_interview_data/image_data'
image_paths = [os.path.join(image_dir, f'train_{i:05}.png') for i in range(num_train+1)]

flattened_images = []
i = 0
for path in image_paths:
    image = imread(path) # BW pngs, no grayscale
    flattened_images.append(image.flatten()) # convert it into a single vector
    i += 1
    if i % 1000 == 0:
        print(f'Processed {i} images')
train_array = np.array(flattened_images)
print(train_array.shape)
train_array = np.save('train_array.npy', train_array)
```

Fig. 1. Loading train images

array is the index of corresponding image.

```
train_labels = []
with open('cig_interview_data/train.txt', 'r') as f:
    for line in f:
        label = int(line.strip().split()[-1])
        train_labels.append(label)
train_labels = np.array(train_labels)
print(train_labels)
print(train_labels.shape)
np.save('train_labels.npy', train_labels)
```

Fig. 2. Loading train labels

C. Data Normalization

Since BW images has color range from 0 to 255, we use normalization to scale the color range to 0 to 1. Initially, we perform an additional normalization using the mean and standard deviation. But later we find out that this step has little effect on the accuracy, so we remove it.

III. LOGISTIC REGRESSION

A. Feature Extraction

We perform PCA to extract original $28 * 28$ features into 50. First, we implement the PCA on our own. However, later we find out that sklearn's PCA function is much faster and more efficient. Our own PCA is incapable for handling this large amount of data, particularly on a personal laptop. Also, we use PCA several times in model tuning. Therefore, we use sklearn's PCA function to perform PCA.

```
def pca(X, k):
    #mean of each feature
    (n_samples, n_features) = X.shape
    mean = np.array([np.mean(X[:,i]) for i in range(n_features)])
    #normalization
    norm_X = X - mean
    #scatter matrix
    scatter_matrix = np.dot(np.transpose(norm_X), norm_X)
    #Calculate the eigenvectors and eigenvalues
    eig_val, eig_vec = np.linalg.eig(scatter_matrix)
    eig_pairs = [(np.abs(eig_val[i]), eig_vec[:,i]) for i in range(n_features)]
    # sort eig_vec based on eig_val from highest to lowest
    eig_pairs.sort(reverse=True)
    # select the top k eig_vec
    feature = np.array([ele[1] for ele in eig_pairs[:k]])
    #get new data
    data = np.dot(norm_X, np.transpose(feature))
    return data
```

Fig. 3. Our PCA

B. Code

We implement the logistic regression basing only on numpy. Also, we use the following accuracy() to evaluate the model. It's worth mentioning that to fit this model, we need to first


```
def trainPerceptron(train_set, train_labels, max_iter):
    #Write code for Mp4
    lr = 0.01
    n_samples, n_features = train_set.shape
    W = np.zeros(n_features)
    b = 0
    for i in range(max_iter):
        for j in range(n_samples):
            y_pred = np.dot(W, train_set[j]) + b
            # gradient
            if y_pred <= 0: y_pred = 0
            else: y_pred = 1
            W += lr * (train_labels[j] - y_pred) * train_set[j]
            b += lr * (train_labels[j] - y_pred)
            if j%100 == 0 and i %10 ==0: print("Iteration: ", i, "Sample: ", j)
        return W, b

def classifyPerceptron(train_set, train_labels, dev_set, max_iter):
    #Write code for Mp4
    W, b = trainPerceptron(train_set, train_labels, max_iter)
    n_samples, n_features = dev_set.shape
    y_pred = np.zeros(n_samples)
    ret = []
    y_pred = np.dot(dev_set, W) + b
    for i in range(len(y_pred)):
        if y_pred[i] <= 0: ret.append(0)
        else: ret.append(1)
    return ret
```

Fig. 9. Perceptron

fact that blurred and unblurred images are quite different, and thus a local minimum of loss function is easily reached. This also lead us to wonder if there is a even simpler way to solve the problem.

V. A SIMPLE APPROACH

When I try to improve the logistic regression, it's really hard due to the discussion above. However, I find some interesting feature. If you compare a blurred and an unblurred image of the same digit, you will realize that the most significant difference is the number of grey pixels. Typically, a blurred image will have more pixels with value ranging between $(0.3 * 255, 0.7 * 255)$, about 2 to 6 times of the number of pixels in this range for an unblurred image. Therefore, we come up with a simple algorithm that counts the number of pixels in range $(0.3 * 255, 0.7 * 255)$, use different thresholds values to categorize whether a picture is blurred or not.

A. Code

First, process the data to get an array to count the number of pixels that lies in the range for each image. And make the decision completely on the thresholds, and evaluate the accuracy.

B. Threshold

We try several threshold, and find the best one on both train set and validation set is 100. The accuracy can achieve 0.89 on both train set and validation set.

VI. CNN

A. Data Preprocessing

The reason we adopt CNN to solve the problem is that CNN is powerful in handling image data, especially in this case the image is in good quality and the output is only binary. Its convolution layer will be so happy to get this problem done. Here we directly load all the train images into a $(28*28, 40000)$ array. Given the dicoverly that both blurred and unblurred images have same value in black area, one may

```
Iteration: 0 Sample: 0
Iteration: 0 Sample: 10000
Iteration: 0 Sample: 20000
Iteration: 0 Sample: 30000
Iteration: 10 Sample: 0
Iteration: 10 Sample: 10000
Iteration: 10 Sample: 20000
Iteration: 10 Sample: 30000
Iteration: 20 Sample: 0
Iteration: 20 Sample: 10000
Iteration: 20 Sample: 20000
Iteration: 20 Sample: 30000
Iteration: 30 Sample: 0
Iteration: 30 Sample: 10000
Iteration: 30 Sample: 20000
Iteration: 30 Sample: 30000
Iteration: 40 Sample: 0
Iteration: 40 Sample: 10000
Iteration: 40 Sample: 20000
Iteration: 40 Sample: 30000
Iteration: 50 Sample: 0
Iteration: 50 Sample: 10000
Iteration: 50 Sample: 20000
Iteration: 50 Sample: 30000
Iteration: 60 Sample: 0
...
Iteration: 90 Sample: 20000
Iteration: 90 Sample: 30000
accuracy on val set: 0.8932
accuracy on train set: 0.89905
```

Fig. 10. Accuracy of Perceptron

leave the outer black edges to reduce input size. However, later we find that CNN is so powerful, that the time to train the model in original size is even faster than a for loop to remove all the edge and then train. Therefore, we directly use original size to train the model.

B. Model Construction

We use tensorflow and keras to contruct the model. More specifically, there are three convolutional layers with ReLU activation function, each layer is followed by a max pooling layer. And then we use a flatten layer to convert the 2D array into 1D array, applying a dense layer of 512 neurons to keep information. Finally, a dense layer of 1 neuron is used for binary output.

```

max = 0.7
min = 0.3

train = np.load('train_array.npy')
train = train/255
val = np.load('val_array.npy')
val = val/255
train_extracted = []
val_extracted = []

for i in range(len(train)):
    count = 0
    for j in range(len(train[i])):
        if train[i][j] > min and train[i][j] < max:
            count += 1
    train_extracted.append(count)
for i in range(len(val)):
    count = 0
    for j in range(len(val[i])):
        if val[i][j] > min and val[i][j] < max:
            count += 1
    val_extracted.append(count)
train_extracted = np.array(train_extracted)
val_extracted = np.array(val_extracted)

```

Fig. 11. Simple Algorithm Data

```

max = 0.7
min = 0.3

train = np.load('train_array.npy')
train = train/255
val = np.load('val_array.npy')
val = val/255
train_extracted = []
val_extracted = []

for i in range(len(train)):
    count = 0
    for j in range(len(train[i])):
        if train[i][j] > min and train[i][j] < max:
            count += 1
    train_extracted.append(count)
for i in range(len(val)):
    count = 0
    for j in range(len(val[i])):
        if val[i][j] > min and val[i][j] < max:
            count += 1
    val_extracted.append(count)
train_extracted = np.array(train_extracted)
val_extracted = np.array(val_extracted)

```

Fig. 12. Simple Algorithm




Fig. 13. Thresholds Evaluation

```

# normalization
train_array = train_array / 255.0
val_array = val_array / 255.0
test_array = test_array / 255.0

# reshape
train_array = train_array.reshape((train_array.shape[0], 28, 28, 1))
val_array = val_array.reshape((val_array.shape[0], 28, 28, 1))
test_array = test_array.reshape((test_array.shape[0], 28, 28, 1))

# CNN
model = models.Sequential([
    layers.Conv2D(16, (3,3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(32, (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    # layers.Conv2D(64, (3,3), activation='relu'),
    # layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer = 'adam', loss = 'binary_crossentropy',
              metrics = 'accuracy')

model.fit(train_array, train_labels, epochs = 5, batch_size = 50)

test_loss, test_acc = model.evaluate(val_array, val_labels)

print('Test accuracy:', test_acc)

```

Fig. 14. CNN

VII. TEST SET

We use both logistic regression and CNN to predict the test set and make several prediction files: test_PERCP.txt, test_CNN.txt, test_LR.txt, and test_SIMPLE.txt, with the similar format as the train.txt and validation.txt.

ACKNOWLEDGMENT

Thanks Professor Do and Qian Jiang for the chance to perform this task. To be honest, I'm not sure whether my logistic

C. Output Evaluation

Initially, we set epochs to 5 and batch size to 50. This small epochs completely comes from my faith in CNN. It's still surprising to see this combination of epochs and batch size gives an accuracy of 1. To further concrete this is not a lie, I print the first 31 labels on train set to manually check, and they are indeed the same.

Good work CNN!

```

Epoch 1/5
800/800 [=====] - 4s 5ms/step - loss: 0.0200 - accuracy: 0.9897
Epoch 2/5
800/800 [=====] - 4s 5ms/step - loss: 3.6166e-06 - accuracy: 1.0000
Epoch 3/5
800/800 [=====] - 4s 5ms/step - loss: 1.8019e-06 - accuracy: 1.0000
Epoch 4/5
800/800 [=====] - 4s 5ms/step - loss: 8.1356e-07 - accuracy: 1.0000
Epoch 5/5
800/800 [=====] - 4s 5ms/step - loss: 4.2896e-07 - accuracy: 1.0000
313/313 [=====] - 1s 2ms/step - loss: 2.7407e-07 - accuracy: 1.0000
Test accuracy: 1.0

```

Fig. 15. Output

```
print(train_labels[:30])
print(model.predict(train_array[:30]))
✓ 0.3s

[0 0 1 1 0 0 1 1 0 0 0 1 0 1 0 0 0 0 0 0 0 1 1 0 1 0 1 1 1 1]
1/1 [=====] - 0s 56ms/step
[[1.86145535e-08]
 [1.22099575e-08]
 [1.00000000e+00]
 [1.00000000e+00]
 [4.73422865e-08]
 [6.20106810e-09]
 [1.00000000e+00]
 [1.00000000e+00]
 [7.87131749e-10]
 [1.12355174e-08]
 [3.03961940e-08]
 [1.00000000e+00]
 [1.88784512e-07]
 [1.00000000e+00]
 [3.27918859e-10]
 [1.51878883e-08]
 [4.15277341e-08]
 [2.79649839e-07]
 [1.37465230e-08]
 [2.19778613e-08]
 [9.07751296e-09]
 [1.00000000e+00]
 [1.00000000e+00]
 ...
 [1.00000000e+00]
 [1.00000000e+00]
 [1.00000000e+00]
 [1.00000000e+00]]
```

Fig. 16. manually check

regression is correctly implemented because the accuracy is so low. But this is a good chance to dig deeply into the concept, as previously I will directly import libraries to do the job. I have learned a lot from this implementation.

REFERENCES

- [1] Verma, M. (2022, May 8). Binary classification using Convolution Neural Network (CNN) model. Medium. <https://medium.com/@mayankverma05032001/binary-classification-using-convolution-neural-network-cnn-model-6e35cdf5bdbb>